For a **Computer Vision (CV) Engineer**, the focus shifts from general data science to **spatial data, image processing pipelines, and edge deployment**. Recruiters at this level want to see that you understand the "physics" of an image—lighting, augmentation, and resolution—as much as the neural network architecture.

[Your Full Name]

Computer Vision Engineer | Deep Learning & Perception

[City, State] • [Phone Number] • [Email Address]

[LinkedIn Profile URL] • [GitHub/Portfolio/Model Zoo URL]

Professional Summary

Specialized Computer Vision Engineer with 5 years of experience building scalable perception systems. Expert in **Object Detection, Semantic Segmentation, and Pose Estimation**. Proven ability to bridge the gap between research and production by deploying optimized models on constrained hardware (NVIDIA Jetson, OAK-D). Strong background in traditional OpenCV techniques paired with modern Deep Learning architectures.

Core Technical Skills

Category	Technologies
CV Frameworks	MMClassification, Detectron2, YOLO (v8-v11), MediaPipe
Deep Learning	PyTorch, TorchVision, TensorFlow, Keras
Image Processing	OpenCV, Scikit-Image, PIL, Albumentations
Edge & Deployment	TensorRT, OpenVINO, CoreML, Docker, Triton Inference Server
Tools & Cloud	Label Studio, CVAT, Weights & Biases, AWS (S3, Lambda)

Professional Experience

Intermediate Computer Vision Engineer | [Current Company Name]

March 2023 – Present

- **Developed a real-time multi-object tracking (MOT) system** using ByteTrack and YOLOv10 for urban traffic monitoring, maintaining 30 FPS on edge devices.
- **Designed a custom synthetic data generation pipeline** using Unity Perception/Blender, reducing reliance on manual labeling by 60% while improving model robustness in low-light conditions.
- **Implemented an automated QC (Quality Control) system** using Siamese Networks for anomaly detection in high-speed manufacturing lines, achieving a 99.2% recall rate.
- **Reduced model size by 75%** through pruning and INT8 quantization using NVIDIA TensorRT, enabling deployment on low-power IoT sensors.

Junior Computer Vision Engineer | [Previous Company Name]

June 2021 – February 2023

- **Built a facial recognition and liveness detection module** for a mobile banking app, utilizing MTCNN and custom FaceNet architectures.
- **Optimized image preprocessing kernels** in C++ and OpenCV, reducing end-to-end latency of the vision pipeline by 40ms per frame.
- **Managed a team of 10+ data annotators**, establishing rigorous QA protocols and active learning loops to improve training set diversity.
- **Integrated vision models into ROS (Robot Operating System)** for an autonomous warehouse robot project, focusing on obstacle avoidance and SLAM.

Technical Projects

3D Reconstruction & Point Cloud Processing

- Developed a pipeline to convert 2D stereo images into 3D point clouds using Open3D.
- Implemented ICP (Iterative Closest Point) algorithms for point cloud registration with <2cm error margin.

Medical Imaging: Cell Segmentation

- Fine-tuned a U-Net architecture for identifying cancerous cells in histopathology slides.
- Leveraged **Dice Loss** and **Focal Loss** to handle extreme class imbalance in microscopic imagery.

Education

B.S./M.S. in Computer Science or Electrical Engineering

University Name | Key Coursework: Digital Image Processing, Pattern Recognition, Robotics.

Key Performance Metrics (Add these to your bullets!)

In Computer Vision, numbers matter. When writing your own CV, try to include:

- **mAP (mean Average Precision):** "Increased mAP from 0.72 to 0.85."
- **FPS (Frames Per Second):** "Achieved real-time performance at 60 FPS."
- **Latency:** "Reduced inference time from 200ms to 45ms."
- **Dataset Size:** "Trained on a curated dataset of 2M+ labeled images."