[Your Full Name]

Senior AI Engineer | Machine Learning & MLOps

[City, State] • [Phone Number] • [Email Address]

[LinkedIn Profile URL] • [GitHub/Portfolio URL]

Professional Summary

Results-driven AI Engineer with 4+ years of experience designing, training, and deploying production-grade machine learning models. Expert in Natural Language Processing (NLP) and Computer Vision, with a strong focus on **MLOps** and scalable cloud architecture. Proven track record of reducing model inference latency by 30% and improving predictive accuracy in enterprise environments.

Core Technical Skills

Category	Technologies
Languages	Python (Expert), SQL, C++, Bash
Frameworks	PyTorch, TensorFlow, JAX, Scikit-learn, Hugging Face
AI/ML Domains	LLMs (Fine-tuning, RAG), Computer Vision, Reinforcement Learning
Data & Cloud	AWS (SageMaker, Lambda), GCP (Vertex AI), PostgreSQL, MongoDB
DevOps/MLOps	Docker, Kubernetes, MLflow, DVC, CI/CD, Weights & Biases

---

Professional Experience

Mid-Level AI Engineer | [Current Company Name]

January 2023 – Present

- **Led the development of a Retrieval-Augmented Generation (RAG) pipeline** using LangChain and Pinecone, reducing customer support response times by 45%.

- **Optimized large-scale transformer models** for production deployment using ONNX and TensorRT, achieving a 2.5x speedup in inference.
- **Architected an automated MLOps pipeline** using Kubeflow, reducing the model deployment cycle from weeks to days.
- **Mentored 3 Junior Engineers** on best practices for data versioning and unit testing for ML code.

Machine Learning Engineer | [Previous Company Name]

June 2020 – December 2022

- **Developed and deployed a Computer Vision system** for real-time defect detection in manufacturing, reaching 98.5% mAP (mean Average Precision).
- **Implemented hyperparameter tuning strategies** (Optuna) that increased model F1-score by 12% across five core product features.
- **Collaborated with Data Engineers** to build robust ETL pipelines processing 1TB+ of daily log data using Apache Spark.
- **Managed GPU resource allocation** within a shared cluster, reducing monthly cloud infrastructure costs by $4,000.

Key Projects

Open-Source Contribution: [Project Name/Link]

- Developed a custom PyTorch implementation of [Specific Architecture, e.g., Vision Transformers] which gained 500+ stars on GitHub.
- Reduced memory footprint by 20% through efficient gradient checkpointing.

End-to-End LLM Application: [Project Name]

- Built a personal finance assistant using Llama 3 and a vector database.
- Integrated a FastAPI backend with a React frontend, deployed via Docker on AWS EC2.

Education

M.S. in Computer Science (Specialization in AI)

University Name | Year of Graduation

B.S. in Software Engineering

University Name | Year of Graduation

Certifications

- AWS Certified Machine Learning – Specialty
- DeepLearning.AI TensorFlow Developer Specialization
- NVIDIA Graduate Fellow (Optional/Honors)